

A decorative graphic on the left side of the slide, consisting of three overlapping circles (blue, pink, and dark blue) partially obscured by a light gray vertical bar.

Modeling Continuous-time Event Data with Temporal Point Processes

Oleksandr Shchur

Technical University of Munich

Berlin Time Series Analysis Meetup – October 12th, 2021

Agenda



Basics



Applications



Old-school TPP models



Neural TPP models



Training TPPs

Temporal point process (TPP)

Probability distribution over variable-length continuous-time event sequences




Hospital visits



Financial transactions



Social media posts




Marked TPPs

Marks – additional features associated with each event, such as

- User ID in the social network
- Magnitude of an earthquake
- Location of a disease outbreak


Marks can be continuous, discrete, vector-valued, ...



Time series vs. TPPs


Time series

- Signal measured at regular intervals



Temporal point process

- We model the arrival times
- The number of events is random




A decorative graphic on the left side of the slide, consisting of three overlapping circles (blue, pink, and dark grey) partially obscured by a light grey vertical bar.

Applications


Application: Prediction

- What will be the type of the next event?
- When will the next event of type ● happen?
- How many events of type ◆ will happen in the future?



Application: Earthquake forecasting


- How many aftershocks of magnitude ≥ 4 do we expect in the next month?




[<https://earthquake.usgs.gov/data/oaf/overview.php>]

Application: Structure discovery

Observed event sequences




Influence structure




[[Linderman, Adams, ICML 2014](#); [Xu, Farajtabar, Zha, ICML 2016](#)]

Application: Anomaly detection

Normal data



Is a new sequence
normal or anomalous?



[[Shchur, Türkmen, Januschowski, Gasthaus, Günnemann, NeurIPS 2021](#)]


A decorative graphic on the left side of the slide, consisting of three overlapping circles (blue, pink, and dark grey) partially obscured by a light grey vertical bar.

Old-school TPPs

How can we describe a TPP?

- A TPP is fully specified by its conditional intensity function

$$\lambda(t|\mathcal{H}_t) = \lim_{dt \rightarrow 0} \frac{\Pr(\text{next event} \in [t, t + dt) | \mathcal{H}_t)}{dt}$$




Homogeneous Poisson process

- Simplest possible model – constant intensity

$$\lambda(t|\mathcal{H}_t) = \mu$$

- Events are independent
- Rate of arrival is constant




Inhomogeneous Poisson process

- Intensity changes over time but is independent of history

$$\lambda(t|\mathcal{H}_t) = g(t)$$

- Captures global trends




Hawkes process

- Intensity increases after each event, then decays to the baseline

$$\lambda(t|\mathcal{H}_t) = \mu + \sum_{t_j \in \mathcal{H}_t} \alpha \exp(-\beta(t - t_j))$$

- Events are clustered ("bursty")




Self-correcting process

- Intensity accumulates over time, drops after each event

$$\lambda(t|\mathcal{H}_t) = \exp\left(\mu t - \sum_{t_j \in \mathcal{H}_t} \alpha\right)$$

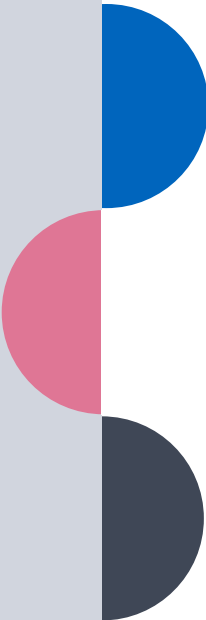
- Events are evenly-spaced



Overview of conventional TPPs

- Conditional intensity $\lambda(t|\mathcal{H}_t)$ fully defines the TPP
- Simple parametric intensity functions
 - ✓ Interpretable
 - ✗ Limited flexibility
- How do we define flexible TPPs that capture complex dependencies?




A decorative graphic on the left side of the slide, consisting of three overlapping circles (blue, pink, and dark grey) partially obscured by a light grey vertical bar.

Neural TPPs: Autoregressive models


TPP as an autoregressive model

- We can equivalently define a TPP by modeling conditional distributions




TPP as an autoregressive model

- We can equivalently define a TPP by modeling conditional distributions



TPP as an autoregressive model

- We can equivalently define a TPP by modeling conditional distributions




Conditional PDF vs. conditional intensity

- Conditional PDF

$$p(t|t_1, \dots, t_{i-1})dt = \Pr(\text{next event} \in [t, t + \Delta t) | \{t_1, \dots, t_{i-1}\})$$


- Conditional intensity

$$\lambda(t|\mathcal{H}_t)dt = \Pr(\text{next event} \in [t, t + \Delta t) | \{\text{no event in } (t_{i-1}, t)\} \cup \{t_1, \dots, t_{i-1}\})$$



Autoregressive neural TPPs


- Main idea: Model the conditional PDF with neural networks



[Du, Dai, Trivedi, Upadhyay, Gomez-Rodriguez, Song, KDD 2016; Shchur, Biloš, Günnemann, ICLR 2021]

Encoding the history into a vector

1. Representing events as feature vectors \mathbf{y}_i
 - Inter-event time as feature
 - Continuous-time positional encoding
2. Aggregate feature vectors $\{\mathbf{y}_1, \dots, \mathbf{y}_{i-1}\}$ into a history embedding \mathbf{c}_i
 - RNN
 - Transformers




Modeling the conditional distribution

1. Pick a parametric PDF $p(\cdot | \boldsymbol{\theta})$ over $[0, \infty)$
 - Simple distribution (Gamma, log-normal, ...)
 - Mixture distribution
 - Normalizing flows

2. Compute parameters from the history embedding

$$\boldsymbol{\theta}_i = \sigma(\mathbf{W}\mathbf{c}_i + \mathbf{b})$$

3. Obtain the conditional distribution
$$p(t_i | t_1, \dots, t_i) = p(t_i - t_{i-1} | \boldsymbol{\theta}_i)$$




A decorative graphic on the left side of the slide, consisting of three overlapping circles (blue, pink, and dark grey) partially obscured by a light grey vertical bar.

Neural TPPs: Continuous-time state evolution

Continuous-time state evolution

- State $\mathbf{h}(t)$ evolves in continuous time
- State directly defines the intensity, e.g., $\lambda(t|\mathcal{H}_t) = \exp(\mathbf{w}^T \mathbf{h}(t))$



[[Mei, Eisner, NeurIPS 2017](#); [Jia, Benson, NeurIPS 2019](#); [Rubanova, Chen, Duvenaud, NeurIPS 2019](#)]

State evolution

- 1 Start with initial state $\mathbf{h}(0) \in \mathbb{R}^D$
- 2 State evolves continuously between events

$$\mathbf{h}(t + \Delta t) = \text{Evolve}(\mathbf{h}(t), t, t + \Delta t)$$

e.g., neural ODE


$$\mathbf{h}(t + \Delta t) = \mathbf{h}(t) + \int_t^{t+\Delta t} \frac{\partial \mathbf{h}(u)}{\partial u} du$$

- 3 Discrete update after each event

$$\mathbf{h}(t_i^+) = \text{Update}(\mathbf{h}(t_i), \mathbf{y}_i)$$

e.g., RNN update

$$\mathbf{h}(t_i^+) = \tanh(\mathbf{W}\mathbf{h}(t_i^+) + \mathbf{V}\mathbf{y}_i + \mathbf{b})$$



Autoregressive vs. continuous-time TPPs

- Autoregressive
 - ✓ Closed-form likelihood evaluation
 - ✓ Closed-form sampling
- Continuous-time state evolution
 - ✓ Naturally handle missing data
 - ✗ Require numerical integration for sampling and training


A decorative graphic on the left side of the slide, consisting of three overlapping circles (blue, pink, and dark grey) partially obscured by a light grey vertical bar.

Parameter estimation

Learning: Maximum likelihood (MLE)

- Log-likelihood function

$$\log p_{\theta}(\{t_1, \dots, t_N\}) = \underbrace{\sum_{i=1}^N \log \lambda(t_i | \mathcal{H}_{t_i})}_{\text{Probability of events at times } \{t_1, \dots, t_N\}} - \underbrace{\int_0^T \lambda(u | \mathcal{H}_u) du}_{\text{Probability of NO events in the rest of the interval}}$$



Learning: Sampling-based losses

- Compute loss based on sampled trajectories

$$\max_{\theta} \mathbb{E}_{\{t_1, \dots, t_N\} \sim \text{TPP}_{\theta}} [f(\{t_1, \dots, t_N\})]$$

- Equivalent to the reparametrization trick – but now for TPPs

Application	TPP_{θ}	$f(\{t_1, \dots, t_N\})$
Generative modeling	Learned model	Sample quality
Reinforcement learning	Policy	Reward function
Variational inference	Approximate posterior	Evidence lower bound

[[Yan, Liu, Shi, Li, Zha, IJCAI 2018](#); [Upadhyay, De, Gomez-Rodriguez, NeurIPS 2018](#); [Shchur, Gao, Biloš, Günnemann, NeurIPS 2020](#)]

Summary

- TPPs – probabilistic models for continuous-time event data
- Two equivalent ways to define a TPP: conditional intensity or conditional PDFs
- Neural TPPs – flexible alternatives to conventional models
- Lots of existing applications – many more to discover



Oleksandr Shchur
TU München
shchur.github.io